

Welcome to the Workshop!

Before we start:

Please open up Claude dot AI (or LLM of choice)

Please also open this repo:

<https://github.com/buddytrip05/specs-and-context>



SCAN ME

AI PM Workshop

From Specs to Context

How PMs learn and ship faster in an AI-Driven SDLC

Buddy Triplett
Product Tank

90 min

hands-on

5

Examples

1

functional file system

About me

I'm Buddy

Fast Facts:

- Principal PM, SpecBooks
- Former PM at **DraftKings** and **General Motors**
- Been doing product for 10 years
- Prior to Product, I led nuclear power plant operations on submarines in the Navy
- In my husband and girl dad era
- Loves gardening - best not to get me talking too much about it or we won't finish tonight
- I like to run for fun

My Why



The only thing to remember today

**LLM output quality improves when PMs
improve the system around the model —
not just their prompting.**

We're going to prove that claim five times today.

Write that down for the test at the end.

SCENE SETTING

Meet Maya.

Senior PM, Terracart

Owens the checkout experience

Has been doing this for 6 years



Terracart is a fictional mid-market outdoor e-commerce platform.

*Monday morning coffee and checkout
funnel dashboard review*

Cart → checkout	84%
Checkout → purchase	32%

SCENE SETTING

Meet Maya.

Senior PM, Terracart

Owens the checkout experience

Has been doing this for 6 years



Terracart is a fictional mid-market outdoor e-commerce platform.

*Monday morning coffee and checkout
funnel dashboard review*

Cart → checkout	84%
-----------------	-----


Checkout → purchase	32%
---------------------	-----

- Her peer said 54% conversion
- She blinks. She refreshes.
- Still 32%

Something is very wrong.

THE FUNNEL DATA


The leakage is hiding in the aggregate.



Nobody drilled into the step-level data.

How Maya found the problem

Three days. One \$2 million discovery.



The data wasn't missing. The context was missing. A benchmark from a casual meetup conversation is what unlocked the question.

Three failure modes.

01

Mobile form UX

Keyboard occludes the expiration field on iOS.

58% abandonment on mobile vs. 31% on desktop.

02

Error messaging

Raw API errors passed verbatim.

"Transaction could not be completed. Please try again."
— that's it.

03

Confirmation email

Silent failure.

78% of support tickets are orders that processed fine.

97 duplicate orders this quarter.

This is not a demand problem. It is a plumbing problem.

Maya's situation right now

1 Day

Prioritize approach and hand-off a draft PRD for engineering review (Thursday EOD)

2 Days

Exec summary and alignment review with leadership (Friday)

5 Days

Review tech plan for start of sprint (Monday)

*Data she has: chaotic meeting notes · 5 interview transcripts · 20 support tickets ·
and a weak PRD draft she wrote at 11pm Monday night*

What we're building today

01

Notes → PRD
draft

15 min

02

Transcripts →
insight memo

15 min

03

Insight memo → 3
audiences

10 min

04

Context folder

15 min

05

Establish PRD
reviewer agents

10 min

*The one that changes
how you work*

01

Exercise 1

Messy notes

→ PRD draft

Files needed:

meeting-notes-raw.md

prd-template.md

15 min

EXERCISE 1

Do the bad prompt first.

Step 1 — The bad prompt (run this first)

"Write a PRD for fixing Terracart's checkout abandonment problem."

What did Claude have to make up? Where is it confidently wrong?

Step 2 — The structured prompt

Paste: meeting-notes-raw.md + prd-template.md into Claude with the structured prompt from [04-prompts/prompt-library.md](#)

Constraint: mark any unknown as [UNKNOWN — needs resolution]. No invention.

Step 3 — Pick the weakest section. Follow up to sharpen it.

02

Exercise 2

Interview transcript → insight memo

Files needed:

customer-interview-transcript.md

support-tickets-sample.md

15 min

EXERCISE 2

Sourcing rules change what AI does with your content.

Without sourcing rules

AI summarizes. Generic themes. Phrases like "users feel frustrated." No evidence. No quotes. Nothing you can defend in a meeting.

With sourcing rules

AI retrieves. Real quotes, labeled by source. Evidence table. Themes you can take into a room and defend.

The rule: "Do not invent or paraphrase quotes. Use exact language from the sources."

03

Exercise 3

One source → three audiences

Exec brief — 5 bullets, dollar-forward

Engineering handoff — scope, blockers, open Qs

Slack update — 3 sentences, human

10 min

EXERCISE 3

The source-of-truth effect.

VP + CEO

They want:

The problem, the impact, risk, timeline, solution direction.

Format:

**5 bullets. ≤20 words each.
Dollar-first.**

Lead engineer

They want:

Scope, blockers, open questions, your ask from them.

Format:

**2-sentence context → workstreams
→ open Qs → what you need by
Monday.**

Product team (Slack)

They want:

What, why, what's next — read in 30 seconds.

Format:

3–5 sentences. No bullets. No "hey team!"

Notice how your Exercise 2 output affects Exercise 3. Better source = better output. Every time.

04

Exercise 4

Build your context folder

The one that makes every future conversation better.

The one that makes you go faster.

This one you take back to your actual job.

15 min

EXERCISE 4

The context folder.

[product]-context/

CONTEXT.md

/research/

interview-transcripts/

support-data/

/specs/

active-prd/

/decisions/

decision-log.md

/stakeholders/

audience-guide.md

CONTEXT.md can contain:

- Company/Product overview
- Current KPIs and company goals
- User journeys and jobs to be done
- Users/personas
- Feature Requirements
- Constraints and non-goals
- Definitions
- Open questions
- Team structure and roles/personas
- Voice and writing styles
- Areas to avoid

Paste CONTEXT.md into any Claude conversation. Or reference from Claude Code. Watch the output quality difference immediately.

05

Exercise 5

PRD reviewer workflow

Generate, then audit.

Leverage AI as a cross-functional stress test.

Files needed: weak-prd-draft.md

10 min

EXERCISE 5

Run the same PRD through three sets of AI eyes.

VP of Product

Cares about: Business impact, timeline, risk. Are we solving the right problem? What's the ask?

Customer support lead

Cares about: Will the fix reduce ticket volume? Is there a comms plan? Have edge cases been thought through?

Skeptical engineer

Cares about: Specificity. Missing constraints. Technical assumptions. Are success metrics actually measurable?

Then: turn the review pattern into a reusable checklist for any PRD you write.

In Closing

- | | | |
|----|---------------------------|--|
| 01 | Notes → PRD | <i>Template anchoring + sourcing constraints</i> |
| 02 | Transcript → insight memo | <i>Sourcing rules = defensible evidence</i> |
| 03 | One source → 3 audiences | <i>Audience framing = format transformation</i> |
| 04 | Context folder | <i>Systems thinking > prompt cleverness</i> |
| 05 | PRD reviewer | <i>Generate, then audit. Always.</i> |

The model does the drafting. The PM does the thinking.

One thing to do Monday: use one of today's prompts on a real artifact you have to ship.

Lenny's Newsletter Meetup



SCAN ME



Lenny's Newsletter Meetup in Dallas (with free drinks courtesy of Clerk)

APR 22 **Wednesday, April 22**
6:00 PM - 8:00 PM

📍 **Chicken N Pickle - Grapevine** ↗
Grapevine, Texas

Starting in 14d 23h

You're In

📅 Add to Calendar 🔗 📧

No longer able to attend? Notify the host by canceling your registration.

Get Ready for the Event
Profile Complete · Reminder: Email ➔

You have manage access for this event. [Manage](#) ↗

Presented by
Lenny's Newsletter ... > [Subscribe](#)

All meetups are hosted by members of the Friends of Lenny's Newsletter community! Learn more at lennysnewsletter.com/community 🍷

📄 📱 📧 📧

Thank you!

Let's Connect on LinkedIn



SCAN ME

Appendix

Advanced tips for AI-native PMs

Three techniques for PMs who already have a workflow (plus a bonus)
and want to make it significantly more powerful.

01

Teach Claude your taste

02

Audit before you read

03

Chain, don't restart

Teach Claude your taste, not just your task.

The problem

Most PMs use Claude transactionally — describe the task, get an output, edit it to sound like themselves. You're re-teaching your voice from scratch every single conversation.

The technique

Paste 3–4 pieces of writing you're proud of. Ask Claude to identify your patterns and write a style guide you can reuse in CONTEXT.md.

```
Analyze these examples of my writing.  
Identify the specific patterns –  
sentence structure, level of detail,  
what I include vs. omit, how I handle  
uncertainty, how I open and close.
```

```
Write a style guide I can paste into  
future conversations to reproduce  
this voice.
```

The delta between "Claude's version" and "my version" often collapses. You edit for substance, not style.

Use Claude to audit its own output before you do.

The problem

Most people treat Claude's first output as the thing to edit instead of reading the whole document looking for problems. That's the slow, unfocused way.

The technique

Send a second message immediately after generation, before you read the output carefully. Ask Claude to surface its own failure modes first.

```
Before I review this – audit it yourself.
```

- What did you have to assume or invent because I didn't give you enough info?
- Where are you least confident in the accuracy or specificity?
- What would a skeptical engineer / VP / customer push back on first?
- What's the weakest sentence?

You're not reading the whole doc looking for problems. You go straight to the flagged sections. Works on anything.

Chain outputs, don't restart conversations.

The problem

Intermediate users start a new Claude conversation for every task. No shared context. Every output is an island. The exec brief can contradict the PRD.

The technique

Explicitly hand each output forward as the authoritative source for the next task in the same conversation (context window). Build a pipeline, not a pile of tabs.

```
Use this memo as the authoritative
source for everything that follows.
Do not introduce information not
present in it.
```

```
First task: write the exec brief.
After I approve it → eng handoff.
After that → Slack update.
```

One session. Exec brief, eng handoff, Slack update — all consistent with each other. No more "the brief says X but the PRD says Y."

Bonus concept

The attention budget.

Straight from Anthropic's engineering team.

"Context must be treated as a finite resource with diminishing marginal returns. Every new token introduced depletes the [attention] budget by some amount, increasing the need to carefully curate the tokens available to the LLM."

— Anthropic Engineering, *Effective Context Engineering for AI Agents* (2025)

What this means for PMs

Too little context

Claude guesses. Generic output. Could apply to any company, any product.

Too much context

Attention diluted. The specific thing you needed gets lost in the noise.

Goldilocks

Smallest set of high-signal context that fully defines the task.

The test: would a smart new hire need this to do the task? If no — leave it out.